

Gov 2002: Problem Set 10

Spring 2023

Problem Set Instructions

This problem set is due on April 29, 11:59 pm Eastern time. Please upload a PDF of your solutions to gradescope. We will accept hand-written solutions for problem 1 but we strongly advise you to typeset your answers in Rmarkdown. Problems 2-4 should be typeset. Please list the names of other students you worked with on this problem set.

Question 1

The standard output from OLS will give the standard errors for the estimated coefficients, but often we want to obtain measures of uncertainty for the predicted value of Y_i given some value of X_i (that is, the conditional expectation function). Using the example from lecture, we might be interested in the average wait times to vote for individuals making \$25,000, \$50,000, or \$100,000 in annual income, along with measures of uncertainty around those estimates. In this problem we will look at how to calculate interval estimates for these predicted values. Assume the following *true* population model for $Y_i|X_i$:

$$Y_i = \beta_0 + \beta_1 X_i + u_i,$$

where the X_i are random variables and u_i are i.i.d. random variables with $E[u_i | X_i] = 0$ and $Var(u_i | X_i) = \sigma^2$. Suppose we observe a random sample of n paired observations $\{Y_i, X_i\}$. Assume the Gauss-Markov assumptions hold and that we have a large sample. Our goal is to estimate the predicted value at some value $X_i = x$:

$$\mu(x) = E[Y_i | X_i = x] = \beta_0 + \beta_1 x.$$

(a)

Let $\hat{\beta}_0$ and $\hat{\beta}_1$ be OLS estimators of the regression of Y on X . Use what you know about the unbiasedness of OLS estimates to show that $\hat{\mu}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$ is an unbiased estimator of the population quantity $\mu(x) = E[Y_i | X_i = x]$.

(b)

Find the conditional variance of $\hat{\beta}_0$, $Var(\hat{\beta}_0 | X_1, \dots, X_n)$, using these two facts:

$$\text{Cov}(\bar{Y}, \hat{\beta}_1 | X_1, \dots, X_n) = 0 \quad \text{and} \quad \text{Var}(\hat{\beta}_1 | X_1, \dots, X_n) = \frac{\sigma^2}{\sum_{i=1}^n (X_i - \bar{X})^2}.$$

You answer should be in terms of σ^2 and functions of X_i . (Hint: find a useful expression for $\hat{\beta}_0$)

(c)

Find the covariance of the OLS estimates given our X values, $\text{Cov}(\hat{\beta}_0, \hat{\beta}_1 | X_1, \dots, X_n)$, again in terms of σ^2 and functions of the X_i . (Hint: It's not zero!)

(d)

Using what you found in (b) and (c), find the standard error of $\hat{\mu}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$.

(e)

Assume that we don't know σ^2 and instead construct our estimate of the standard error by plugging in for σ^2 our unbiased estimate s^2 using the residuals.

Give the formula for a large-sample 95% confidence interval estimator for $\mu(x) = E[Y | X = x]$ using what you found above and substituting s^2 for σ^2 . How do we interpret this confidence interval?

Question 2

In this problem, you will write your own version of `lm()`'s core functionality. You will need the `subprime` data to answer part of this question.

- (a) Write a generic function that takes two arguments: a **formula** (such as `income ~ loan.amount`), and **data** (a data frame). You may copy the function from last problem set.

Your function should return a `list()` object, with the following elements:

- **coefficients**: the coefficient estimates (using matrix operations)
- **R.squared**: the R^2 of your model
- **standard.errors**: standard errors of your coefficients (assuming homoskedasticity)
- **t.stats**: t -statistics for your coefficients
- **p.vals**: p -values for your coefficients

Do not use `lm()` in your function. We want you to code the estimates using R's matrix operations, such as `t()` (transpose), `%*%` (multiply matrices), `solve()` (invert), `diag()` (extract diagonal).

- (b) Now, test your function again on the `subprime` data. Run the following regression model `income ~ loan.amount + black + woman` using your function from part (a) to `lm()` and compare the outputs.

- (c) Finally, run your regression model with an interaction term, `income ~ loan.amount + black + woman + black:woman`. Compare the outputs to `lm()`. Does the substantive interpretation of your results change from part (b)?

Question 3

For this problem, we are going to look at data on the 2012 Russian Presidential Election. This election was held a year after the 2011 Duma (Parliamentary) Election that, according to OCSE observers, was considered “slanted in favor of the ruling party” - Vladimir Putin’s United Russia.¹ Observers noted a number of irregularities in the voting process, including evidence of ballot-box stuffing at some polling stations. The elections were followed by a series of protests in major cities against the government denouncing the election fraud. As a consequence, the 2012 elections were subject to greater domestic observation efforts, but were still highly skewed in favor of the ultimate winner - Vladimir Putin. As OCSE observers noted, while voting procedures were relatively well followed, the vote count showed many procedural irregularities at around 1/3 of polling stations.²

- (a) We will first look at the election returns at the sub-regional level (roughly equivalent to county-level). Load the dataset `prezElectionSubRepublicLevel.csv` into R. Create a new variable `putinvote` which is equal to the number of votes for Putin (`Putin`) divided by the number of valid ballots cast (`Number.of.Valid.Ballots`) multiplied by 100 (to scale to a percentage from 0 to 100). Further, create another variable `turnout` which is equal to the number of valid ballots cast (`Number.of.Valid.Ballots`) divided by the number of registered voters (`Number.of.Registered.Voters`).
- (b) Make a scatterplot of Putin’s share of the vote on the percentage turnout in the sub-region. Run a linear regression of Putin’s vote share on percent turnout and overlay the regression line on top of the points. Make sure the regression line is clearly visible. Report your estimated regression coefficients and standard errors in a neatly formatted table and interpret the coefficient on % turnout.

Hint: Since there are a lot of datapoints, in order to reduce clutter, you may want to change how the points look (e.g. set `pch=18` if you’re using base R or `shape` in ggplot) and their size (e.g. set `cex = .5` in base or `size = 0.5` in ggplot). You might also want to change the color of the points (e.g. `col = "darkgrey"` in base or `color = "darkgrey"` in ggplot).

- (c) Take a close look at your scatterplot. Just eyeballing it, do you detect any evidence of nonlinearity here?
- (d) Now, test your intuition by making a plot of the residuals from your regression in B) against the fitted values from that regression. Add a smooth loess regression curve to the plot (using `geom_smooth` in ggplot, or the `scatter.smooth` function in base R). Does there appear to be a pattern in the residuals? What does this suggest?
- (e) How might you specify a model for Putin’s vote share that accounts for the pattern

¹<http://www.osce.org/odihr/elections/86981>

²<http://www.osce.org/odihr/elections/88661>

in (c-d)? Run a regression using that model specification and report your results in a nicely formatted table below.

- (f) Make a plot of the residuals from your regression in (e) against the fitted values from the same regression. Add a smooth loess regression curve to the plot as in (d). Does the pattern from (c-d) remain?

Question 4

In the book, “The Economic Effects of Constitutions,” Torsten Persson and Guido Tabellini examine possible institutional determinants of economic performance. A subset of the data they use (“tabellini.dta”) is available on Ed. Below is the list of the variables.

- *rgdph* - real GDP per capita in constant dollars (chain index) expressed in international prices, base year 1985.
- *polityIV* - score for democracy, ranging from +10 (strongly democratic) to -10 (strongly autocratic).
- *gini_8090* - Gini index on income distribution, computed as the average of two data points: the observation closest to the 1980 and the observation closest to the 1990. When only one of the two years year is available, only that year is included.
- *trade* - sum of exports and imports of goods and services measured as a share of GDP.

We will use these data to investigate the validity of the homoskedasticity assumption.

- a. Regress GDP on the polity, gini, and trade variables. Based on your substantive knowledge, do you think the homoskedasticity assumption will hold for this model? We are not looking for a specific answer, only that you explain your reasoning.
- b. Take the results of your regression from part (a) and produce a plot with the absolute value of the residuals on the y-axis and the fitted values on the x-axis. Add a lowess line. Make sure to label your graph. This type of plot is called a spread-location plot and is commonly used to assess the homoskedasticity assumption. What do you notice? Do you think the errors are homoskedastic?
- c. Calculate the heteroskedasticity robust variance-covariance matrix by hand-coding the matrix algebra. Then, check your answer using the “sandwich” package. Make sure to use the small sample correction.
- d. Take the vcov matrix from part (c) and calculate the heteroskedastic robust standard errors for your regression from part (a) and compare to the results of the original regression. What do you notice? Include regression tables with and without heteroskedastic robust standard errors in your answer.